

# CS301: HIGH PERFORMANCE COMPUTING

## ASSIGNMENT TWO

**Submitted By:**

**Shaleen Kumar Gupta (201301429)**

**Visharad Bansal (201301438)**

**Q1.**

### **PROBLEM STATEMENT**

Calculate the value of pi by sum of series.

### **METHOD**

$$\pi = 4 * [ 1 - (1/3) + (1/5) - (1/7) + \dots ]$$

### **COMPLEXITY**

The serial code runs in  $O(n)$  time, where  $n$  is the number of integration steps.

OpenMP is used to divide the serial code among 4 threads, giving a theoretical speedup of 4x


### **OPTIMIZATION STRATEGY**

2 out of 4 threads (thread 0 and 2) are assigned all the positive terms of the series while thread 1 and 3 sum up all the negative terms. Then a reduction clause is used to calculate the final sum.

### **OBSERVATIONS AND PROBLEMS FACED**

Using the OpenMp reduction clause gives a much better time compared to using an array to store local sums. This is probably because of the load and store overheads associated with using an array. Also, rather than declaring the variables inside each thread, they were declared outside in the serial part and made private to each thread. This also helped improve the time for the parallel code.

However, each step still requires a division to be performed, which is an intensive task. There are probably other methods such as Monte Carlo which will give an even better performance.



Comparison of serial and parallel times

## HARDWARE DETAILS

No. of Processors : 3  
vendor\_id : Genuine Intel  
CPU family : 6  
model : 60  
model name : Intel(R) Core(TM) i5-4590 CPU @ 3.30GHz  
stepping : 3  
CPU MHz : 800.000  
cache size : 6144 KB  
siblings : 4  
CPU cores : 4

## INPUT PARAMETERS


Number of steps (n).

n was taken from 10 to 10,000


## OBSERVATIONS

Both serial and parallel codes gave almost linear speedup with respect to the problem size, with a few kinks, specially for small values of n.

Time was calculated using `omp_get_wtime()` library function include in the “omp.h” header file.



Serial Scaling




Parallel Scaling

## SPEED UP

The maximum speed up achieved was ~2.2

Speed Up was calculated as:

Speed Up= (Time taken by serial code) / (Time taken by parallel code)



Speed Up v/s N

## Q2.

### PROBLEM STATEMENT

Calculate the Nth Fibonacci number.

### METHOD

This is the basic algorithm:

```
fib[0]=fib[1]=1;  
fib[n]=fib[n-1]+fib[n-2]
```

Matrix method:

$$\begin{pmatrix} 0 & 1 \\ 1 & 1 \end{pmatrix}^n = \begin{pmatrix} F_{n-1} & F_n \\ F_n & F_{n+1} \end{pmatrix}$$

### COMPLEXITY

After memoization, the code runs in  $O(n)$  time. An array `fib[]` is used to store the values. The matrix implementation runs in  $O(n)$  time, as the multiplication is for a  $2 \times 2$  matrix.

### OPTIMIZATION STRATEGY

#### IMPLEMENTATION 1

For every call to the function `fibonacci(n)` (which returns the  $n$ th fibonacci number), two threads are spawned: one which calls `fibonacci(n-1)` and the second which calls `fibonacci(n-2)`. Then the threads are made to synchronize via a wait call, and the master adds up the two values and returns the answer.

To reduce the time taken, fibonacci numbers are stored in the array `fib[]` the first time they are calculated. The function then simply checks if `fib[n]` is non zero and returns it.

Also, for  $n < 20$  the parallel code was reduced to serial, in the hope of a better speedup.

#### IMPLEMENTATION 2

Parallel matrix multiplication was used to calculate the  $n$ th fibonacci number as per the


formula described earlier. The outer loop was parallelized and the *#pragma for* construct was used.

## OBSERVATIONS AND PROBLEMS FACED

Spawning of new threads and subsequent parallelization of the the code is useful for computation intensive programs. The parallel code for fibonacci fails to give a speedup because the overhead (the master creating new threads, dividing tasks and synchronization at the end) is much more compared to the task performed by each thread, which is in most cases just to call a function and return a value. As a result, the serial code always gave a better time compared to the parallel code.


There is also the size limitation. The datatype long long was used to store the numbers, and it only stored upto the 45th fibonacci number accurately before overflowing. Maybe with a higher value of n a speedup  $> 1$  will be observed.

Graph 1: Serial Implementation with Recursion




Due to memoization, we observe certain spikes

Graph 2: Parallel Implementation with Recursion



Here we observe a more or less linearly increasing time due to the high number of threads spawned


### Graph 3: Parallel Matrix Implementation




Implementation using the matrix method takes a much longer time compared to recursive implementations. This is because of the intensive matrix multiplications (high number of saxpy operations)



Graph 4: Serial and Parallel Implementation with recursion



The SpeedUp Curve



The speedup is less than 1. It hovers around 1 for  $n < 20$  as both codes follow the serial implementation. For  $n > 20$ , the parallel overheads is too high and brings the speedup to an average of 0.5

## **HARDWARE DETAILS**

No. of Processors : 3  
vendor\_id : Genuine Intel  
CPU family : 6  
model : 60  
model name : Intel(R) Core(TM) i5-4590 CPU @ 3.30GHz  
stepping : 3  
CPU MHz : 800.000  
cache size : 6144 KB  
siblings : 4  
CPU cores : 4

## **INPUT PARAMETERS**

Number of steps (n).  
n was taken upto 45.

**Q3.**

### **PROBLEM STATEMENT**

Write a serial code and a parallel code for performing bubble sort and odd-even sort

### **METHOD**

Bubble sort, is a simple sorting algorithm that repeatedly steps through the list to be sorted, compares each pair of adjacent items and swaps them if they are in the wrong order. The pass through the list is repeated until no swaps are needed, which indicates that the list is sorted. Although the algorithm is simple, it is too slow and impractical for most problems even when compared to insertion sort.

Odd-even sort is a sorting algorithm developed originally for use on parallel processors with local interconnections. It is a comparison sort related to bubble sort, with which it shares many characteristics. It functions by comparing all odd/even indexed pairs of adjacent elements in the list and, if a pair is in the wrong order (the first is larger than the second) the elements are switched. The next step repeats this for even/odd indexed pairs (of adjacent elements). Then it alternates between odd/even and even/odd steps until the list is sorted.

### **THEORETICAL SPEEDUP**

OpenMP is used to divide the serial code among 4 threads, giving a theoretical speedup of 4x

### **OPTIMIZATION STRATEGY**


#### Algorithm for Parallel Bubble Sort

1. For  $k = 0$  to  $n-2$
2. If  $k$  is even then
3.     for  $i = 0$  to  $(n/2)-1$  do in parallel
4.     If  $A[2i] > A[2i+1]$  then
5.     Exchange  $A[2i] \leftrightarrow A[2i+1]$
6. Else
7.     for  $i = 0$  to  $(n/2)-2$  do in parallel
8.     If  $A[2i+1] > A[2i+2]$  then
9.     Exchange  $A[2i+1] \leftrightarrow A[2i+2]$
10. Next  $k$


## COMPLEXITY

Steps 1-10 is a one big loop that is represented  $n - 1$  times. Therefore, the parallel time complexity is  $O(n)$ . If the algorithm, odd-numbered steps need  $(n/2) - 2$  processors and even-numbered steps require  $(n/2) - 1$  processors. Therefore, this needs  $O(n)$  processors.


## OBSERVATIONS AND PROBLEMS FACED




Time taken vs Problem Size for Bubble Sort in Serial Implementation



Time taken vs Problem Size for Bubble Sort in Parallel Implementation



Time taken vs Problem Size for Odd Even Sort in Serial Implementation



Time taken vs Problem Size for Odd Even Sort in Parallel Implementation

## HARDWARE DETAILS


No. of Processors : 3  
vendor\_id : Genuine Intel  
CPU family : 6  
model : 60  
model name : Intel(R) Core(TM) i5-4590 CPU @ 3.30GHz  
stepping : 3  
CPU MHz : 800.000  
cache size : 6144 KB  
siblings : 4  
CPU cores : 4

## INPUT PARAMETERS

Number of steps (n).


n was taken from 100 to 2000 with n being doubled in each iteration.

## OBSERVATIONS



Time taken vs Problem Size for Bubble Sort in Parallel Implementation and Serial Implementation

In the Bubble Sort implementations, we observed that the time taken in Parallel implementation is slightly higher for smaller size arrays while it goes below the Serial curve after a certain point. This depicts the expected outcome. For smaller input sizes, we experience more parallelization overheads than the optimization achieved.




Time taken vs Problem Size for Odd Even Sort in Parallel Implementation and Serial Implementation

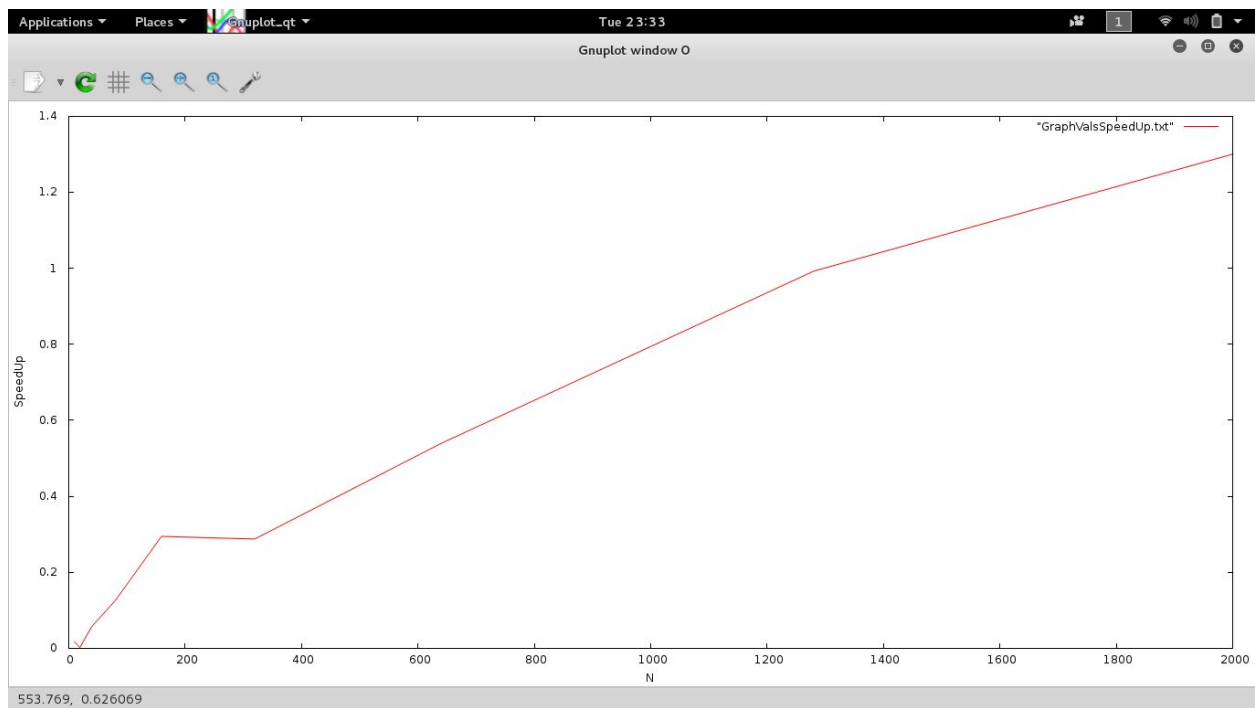
Time was calculated using `omp_get_wtime()` library function include in the "omp.h" header file.

## SPEED UP





Problem Size vs SpeedUp for Bubble Sort in Parallel Implementation against Serial Implementation



Problem Size vs SpeedUp for Bubble Sort in Parallel Implementation against Serial Implementation

Speed Up was calculated as:

Speed Up= (Time taken by serial code) / (Time taken by parallel code)